


CSE 471: MACHINE LEARNING

Deep Learning

Outline

- Feedforward neural network
 - Multi layer perceptron (MLP)
- Convolutional neural network
- Recurrent neural network
-


Perceptron



Multi layer perceptron

Deep learning network

A deep learning network has longer computation paths, allowing each variable to interact with all the others.



Feedforward Network

$$a_j = g_j\left(\sum_i w_{i,j}a_i\right) \equiv g_j(in_j)$$

- g_j – Nonlinear activation function
- There is an intercept b . *Think of it as weight given to a fixed +1 activation node.*

$$a_j = g_j(\mathbf{w}^\top \mathbf{x})$$

Nonlinearity (Sigmoid)

$$\sigma(x) = 1/(1 + e^{-x})$$

Nonlinearity




$$\sigma(x) = 1 / (1 + e^{-x})$$

$$\text{ReLU}(x) = \max(0, x)$$


$$\text{softplus}(x) = \log(1 + e^x)$$

$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1} = 2\sigma(2x) - 1$$


Nonlinearity



(a)





(b)





(c)

Activation functions commonly used in deep learning systems: (a) the logistic or sigmoid function; (b) the ReLU function and the softplus function; (c) the tanh function.

ReLU




Sigmoid




$$\begin{aligned}\frac{d}{dx} \text{sigmoid}(x) &= \frac{\exp(-x)}{(1 + \exp(-x))^2} \\ &= \text{sigmoid}(x) (1 - \text{sigmoid}(x))\end{aligned}$$

Tanh




$$\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$




$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$

Computation Graph




(a)



(b)

Forward computation




$$\begin{aligned}\hat{y} &= g_5(in_5) = g_5(w_{0,5} + w_{3,5}a_3 + w_{4,5}a_4) \\ &= g_5(w_{0,5} + w_{3,5}g_3(in_3) + w_{4,5}g_4(in_4)) \\ &= g_5(w_{0,5} + w_{3,5}g_3(w_{0,3} + w_{1,3}x_1 + w_{2,3}x_2) \\ &\quad + w_{4,5}g_4(w_{0,4} + w_{1,4}x_1 + w_{2,4}x_2))\end{aligned}$$


$$h_{\mathbf{w}}(\mathbf{x}) = \mathbf{g}^{(2)}(\mathbf{W}^{(2)}\mathbf{g}^{(1)}(\mathbf{W}^{(1)}\mathbf{x}))$$

Gradients

$$\begin{aligned}\frac{\partial}{\partial w_{3,5}} \text{Loss}(h_{\mathbf{w}}) &= \frac{\partial}{\partial w_{3,5}} (y - \hat{y})^2 = -2(y - \hat{y}) \frac{\partial \hat{y}}{\partial w_{3,5}} \\&= -2(y - \hat{y}) \frac{\partial}{\partial w_{3,5}} g_5(\text{in}_5) = -2(y - \hat{y}) g'_5(\text{in}_5) \frac{\partial}{\partial w_{3,5}} \text{in}_5 \\&= -2(y - \hat{y}) g'_5(\text{in}_5) \frac{\partial}{\partial w_{3,5}} (w_{0,5} + w_{3,5}a_3 + w_{4,5}a_4) \\&= -2(y - \hat{y}) g'_5(\text{in}_5) a_3.\end{aligned}$$



Gradients



$$\frac{\partial}{\partial w_{3,5}} \text{Loss}(h_{\mathbf{w}}) = -2(y - \hat{y})g'_5(in_5)a_3$$

If we define $\Delta_5 = 2(\hat{y} - y)g'_5(in_5)$ as a sort of “perceived error” at the point where unit 5 receives its input, then the gradient with respect to $w_{3,5}$ is just $\Delta_5 a_3$. This makes perfect sense: if Δ_5 is positive, that means \hat{y} is too big (recall that g' is always nonnegative); if a_3 is also positive, then increasing $w_{3,5}$ will only make things worse, whereas if a_3 is negative, then increasing $w_{3,5}$ will reduce the error. The magnitude of a_3 also matters: if a_3 is small for this training example, then $w_{3,5}$ didn’t play a major role in producing the error and doesn’t need to be changed much.

Gradients

$$\frac{\partial}{\partial w_{1,3}} \text{Loss}(h_{\mathbf{w}}) = -2(y - \hat{y}) g'_5(in_5) \frac{\partial}{\partial w_{1,3}} (w_{0,5} + w_{3,5}a_3 + w_{4,5}a_4)$$


$$= -2(y - \hat{y}) g'_5(in_5) w_{3,5} \frac{\partial}{\partial w_{1,3}} a_3$$

$$= -2(y - \hat{y}) g'_5(in_5) w_{3,5} \frac{\partial}{\partial w_{1,3}} g_3(in_3)$$


$$= -2(y - \hat{y}) g'_5(in_5) w_{3,5} g'_3(in_3) \frac{\partial}{\partial w_{1,3}} in_3$$

$$= -2(y - \hat{y}) g'_5(in_5) w_{3,5} g'_3(in_3) \frac{\partial}{\partial w_{1,3}} (w_{0,3} + w_{1,3}x_1 + w_{2,3}x_2)$$

$$= -2(y - \hat{y}) g'_5(in_5) w_{3,5} g'_3(in_3) x_1.$$



Gradients




$$\frac{\partial}{\partial w_{1,3}} \text{Loss}(h_{\mathbf{w}}) = -2(y - \hat{y}) g'_5(in_5) w_{3,5} g'_3(in_3) x_1$$

If we also define $\Delta_3 = \Delta_5 w_{3,5} g'_3(in_3)$, then the gradient for $w_{1,3}$ becomes just $\Delta_3 x_1$. Thus, the perceived error at the input to unit 3 is the perceived error at the input to unit 5, multiplied by information along the path from 5 back to 3. This phenomenon is completely general, and gives rise to the term **back-propagation** for the way that the error at the output is passed back through the network.

Home work

- $w_{0,3} = 1, w_{0,4} = -1, w_{0,5} = 1$
- $w_{1,3} = 2, w_{1,4} = 4, w_{2,3} = 3$
- $w_{2,4} = 1, w_{3,5} = 2, w_{4,5} = 3$
- $x_1 = 5, x_2 = -3$
- $y = 1$
- Non-linearity = Sigmoid
- $y_{\text{hat}} = ?$
- Calculate the gradients
 - w.r.t the w's.



Gradients

- Vanishing gradient
 - Derivatives can be very close to 0
 - Changing the weights may have negligible effects.
- Automatic differentiation
 - Systematic application of chain rule.